

# Teaching the Teacher

## Day 4 - Afternoon » From text to features: Natural Language Processing «

---

Anne Kroon

a.c.kroon@uva.nl

@annekroon

July 1, 2021

Afdeling Communicatiewetenschap  
Universiteit van Amsterdam

# Today

From text to features: vectorizers

General idea

Pruning

Teaching Q&A

## From text to features: vectorizers

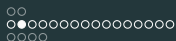
---

From text to features: vectorizers

# From text to features: vectorizers

---

## General idea



# A text as a collections of word

Let us represent a string

```
1 t = "This this is is is a test test test"
```

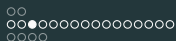
like this:

```
1 from collections import Counter
2 print(Counter(t.split()))
```

```
1 Counter({'is': 3, 'test': 3, 'This': 1, 'this': 1, 'a': 1})
```

Compared to the original string, this representation

- is less repetitive
- preserves word frequencies
- but does *not* preserve word order
- can be interpreted as a vector to calculate with (!!!)



## From vector to matrix

If we do this for multiple texts, we can arrange the vectors in a table.

`t1 = "This this is is is a test test test"`

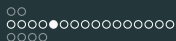
`t2 = "This is an example"`

	a	an	example	is	this	This	test
<i>t1</i>	1	0	0	3	1	1	3
<i>t2</i>	0	1	1	1	0	1	0



*What can you do with such a matrix?  
Why would you want to represent a  
collection of texts in such a way?*





# What is a vectorizer

- Transforms a list of texts into a sparse (!) matrix (of word frequencies)
- Vectorizer needs to be “fitted” to the training data (learn which words (features) exist in the dataset and assign them to columns in the matrix)
- Vectorizer can then be re-used to transform other datasets

## The cell entries: raw counts versus tf·idf scores

- In the example, we entered simple counts (the “term frequency”)



*But are all terms equally important?*

## The cell entries: raw counts versus tf·idf scores

- In the example, we entered simple counts (the “term frequency”)
- But does a word that occurs in almost all documents contain much information?
- And isn't the presence of a word that occurs in very few documents a pretty strong hint?
- **Solution:** Weigh by *the number of documents in which the term occurs at least once* (the “document frequency”)

⇒ we multiply the “term frequency” (tf) by the inverse document frequency (idf)

(usually with some additional logarithmic transformation and normalization applied, see [https://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.TfidfTransformer.html](https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.TfidfTransformer.html))



## tf·idf

$$w_{i,j} = tf_{i,j} \times \log \left( \frac{N}{df_i} \right)$$

$tf_{i,j}$  = number of occurrences of  $i$  in  $j$

$df_i$  = number of documents containing  $i$

$N$  = total number of documents

# Is tf·idf always better?

It depends.

- Ultimately, it's an empirical question which works better (→ machine learning)
- In many scenarios, “discounting” too frequent words and “boosting” rare words makes a lot of sense (most frequent words in a text can be highly un-informative)
- Beauty of raw tf counts, though: interpretability + describes document in itself, not in relation to other documents

## Different vectorizers

1. CountVectorizer (=simple word counts)
2. TfidfVectorizer (word counts (“term frequency”) weighted by number of documents in which the word occurs at all (“inverse document frequency”))

# Internal representations

## Sparse vs dense matrices

- → tens of thousands of columns (terms), and one row per document
- Filling all cells is inefficient *and* can make the matrix too large to fit in memory (!!!)
- Solution: store only non-zero values with their coordinates! (sparse matrix)
- dense matrix (or dataframes) not advisable, only for toy examples



*s p a r s e*

0	7	0	0	0	0	6
0	7	6	3	0	4	0
0	4	3	0	0	0	0
4	2	0	0	0	0	0
0	0	0	0	3	2	4

© Matt Edging

**DENSE**

0	7	0	0	0	0	6
0	7	6	3	0	4	0
0	4	3	0	0	0	0
4	2	0	0	0	0	0
0	0	0	0	3	2	4

<https://matteding.github.io/2019/04/25/sparse-matrices/>

This morning we learned how to tokenize with a list comprehension (and that's often a good idea!). But what if we want to *directly* get a DTM instead of lists of tokens?

# OK, good enough, perfect?

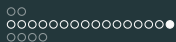
## scikit-learn's CountVectorizer (default settings)

- applies lowercasing
- deals with punctuation etc. itself
- minimum word length > 1
- more technically, tokenizes using this regular expression:  
`r"(?u)\b\w\w+\b"`<sup>1</sup>

```
1 from sklearn.feature_extraction.text import CountVectorizer
2 cv = CountVectorizer()
3 dtm_sparse = cv.fit_transform(docs)
```

---

<sup>1</sup>?u = support unicode, \b = word boundary



# OK, good enough, perfect?

## CountVectorizer supports more

- stopword removal
- custom regular expression
- or even using an external tokenizer
- ngrams instead of unigrams

see

[https://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.CountVectorizer.html](https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html)

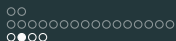
## Best of both worlds

Use the Count vectorizer with a NLTK-based external tokenizer! (see book)

# From text to features: vectorizers

---

Pruning



## General idea

- Idea behind both stopwords removal and tf-idf: too frequent words are uninformative
- (possible) downside stopwords removal: a priori list, does not take empirical frequencies in dataset into account
- (possible) downside tf-idf: does not reduce number of features

Pruning: remove all features (tokens) that occur in less than X or more than X of the documents

## CountVectorizer, only stopwords removal

```
1 from sklearn.feature_extraction.text import CountVectorizer,  
    TfidfVectorizer  
2 myvectorizer = CountVectorizer(stop_words=mystopwords)
```

CountVectorizer, better tokenization, stopwords removal (pay attention that stopwords list uses same tokenization!):

```
1 myvectorizer = CountVectorizer(tokenizer = TreebankWordTokenizer().  
    tokenize, stop_words=mystopwords)
```

Additionally remove words that occur in more than 75% or less than  $n = 2$  documents:

```
1 myvectorizer = CountVectorizer(tokenizer = TreebankWordTokenizer().  
    tokenize, stop_words=mystopwords, max_df=.75, min_df=2)
```

All together: tf-idf, explicit stopwords removal, pruning

```
1 myvectorizer = TfidfVectorizer(tokenizer = TreebankWordTokenizer().  
    tokenize, stop_words=mystopwords, max_df=.75, min_df=2)
```

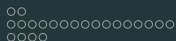


*What is “best”? Which (combination of) techniques to use, and how to decide?*



# Teaching Q&A

---



# NLP and teaching

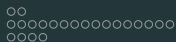
## Teaching experiences

- Transparency
- Students should be able to explain HOW they've preprocessed the data and WHY
- Arguments for preprocessing differ across unsupervised and supervised tasks

# NLP and teaching

## Teaching experiences

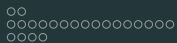
- Rationale for using preprocessing: Why do you use specific techniques?
- For supervised learning: often an empirical question
- Thus: testing different setting and explain what works best. Systematically testing different techniques



# NLP and teaching

## Sentiment analysis

- Dictionary-based approaches
- Keep in mind; what are best practices? Off-the-shelf do not necessarily generalize well.



# Thank you!!

Thank you for your attention!

- Questions? Comments?